CSCI203 Algorithms and Data Structures

String Searching and Improving Sorting II

Lecturer: Dr. Xueqiao Liu

Room 3.117

Email: xueqiao@uow.edu.au

Outline

- String Search
- ▶ Lower Bound for Comparison Sort
 - Decision Tree Model
- Sorting in Linear Time
 - Bucket Sorting
 - Radix Sorting

Looking for Text (in all the right places)

- Consider the problem of String Searching:
 - Given a text, t, is the subtext, s, present in it?
- This problem occurs in many real-life applications:
 - grep;
 - find in a text editor;
 - Genome matching;
 - Google search.

String Search Problem

- Example
 - t = AGCATGCTGCAGTCATGCTTAGGCTA
 - s= GCT
 - s appear three times in t, starting from locations
 6,17,23
- There are a wide number of techniques to achieve this.
- Let us look at a couple of examples.

The Naïve Approach: Linear Search

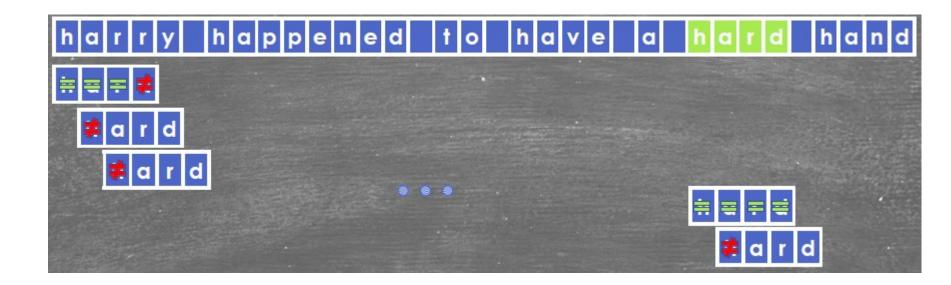
- ▶ The simplest possible approach is linear search:
 - Try to match s starting at each location in t.

```
for i from 0 to length(t) - length(s)
    j=0
    while j < length(s) do
        if (s(j) != t(i+j)) break
        j++
    od
    if j == length(s) print(" string found starting at location " i)
rof</pre>
```

We can see this with an example.

Linear Search: An Example

- \blacktriangleright Let t be the string "harry happened to have a hard hand".
- Let s be the string "hard".
- The search proceeds as follows:



Linear Search ≠ Linear Time Search

- ▶ The outer loop in our algorithm is repeated |t| |s| times.
- Typically the string t is much longer than the string s, so this is $\Theta(t)$.
- The inner loop is repeated up to |s| times for each time round the outer loop.
 - This is $\Theta(s)$
- The total number of comparisons is $\Theta(|s| \times |t|)$.
- ▶ Is this the best we can do?
- The best we can possibly do is $\Theta(|s| + |t|)$;
 - we have to at least look at each string!
- Can we actually achieve this goal of a linear time algorithm?

Linear Time Search

- To do this we will use hashing.
- We compare the hash of string s with the hash of each substring of t with the same length:

harry happened to have a hard hand

Linear Time Search

- This algorithm takes linear time, provided:
 - The hash function only collides rarely;
 - The hash function takes constant time to compute;
 - Independent of the length of string s!
- Surely, the second requirement is impossible.
- ▶ To hash a string of length |s| must take $\Theta(|s|)$ operations.
 - Yes?
 - No!
- Not if we are clever.

Clever Hashing

- We note that we need $\Theta(|s|)$ time to compute h(s).
- We also need $\Theta(|s|)$ time to compute h(t[0..|s|-1), the initial substring of t.
- The trick is to compute the hash of each successive substring of t in constant time.
- If we look closely at these substrings, we see an interesting feature:
 - Successive substrings differ only by two characters.
- ▶ The first character of the first substring; harr
- The last character of the next substring. arry

Rolling Hash

- Maybe we can define a hash function which, given h("harr") can compute h("arry") in constant time.
- Let us define a rolling hash function, r(), so that:
 - h("arry") = r(h("harr"), "h", "y")
- We compute the hash of the next substring by removing the first and appending the new last characters;
- ▶ In this case we remove "h" and append "y".
- If we can compute a rolling hash in constant time then we can do string matching in linear time.
- ► How?

Karp-Rabin String Search

The Karp-Rabin algorithm looks like this:

```
hash_s=hash(s)
hash_t=hash(t[0..length(s)-1])
for i in 0...length(t)-length(s)
   if hash_s == hash_t then
        brute-force compare s and the substring
        if they match print(" string found starting at location
        " i)
   fi
   hash_t=roll(hash_t,t[i],t[i+length(s)])
rof
```

- The function roll(h, p, s) computes the rolling hash of the next substring given the hash of the existing substring, h, with the prefix p removed and the suffix, s, appended.
- We need only find a suitable function roll().

How We Roll

- One popular way to compute roll() is to use something called the Rabin fingerprint.
- We start by treating each symbol in the alphabet as an integer
 use the ASCII code for example.
- We then find a random prime number > the size of the alphabet—let's pick 257.
- We now compute h("harr") as:
 - $257^3 * 104 + 257^2 * 97 + 257^1 * 114 + 257^0 * 114$
 - **1,771,793,837**
- Note: "h" = 104, "a" = 97 and "r" = 114.

The Next Hash

- Given that h("harr") = 1,771,793,837 how do we get h("arry")?
- Simply compute $r(h, p, s) = 257 * (h 257^3 * p) + s$ = 257. (1,771,793,837 - 257³ * 104) + 121
- In this case the result is 1,654,094,526 which is exactly the same as h("arry")

$$257^3 * 97 + 257^2 * 114 + 257^1 * 114 + 257^0 * 121$$

Note: if these values become too large, we can reduce them modulo m, where m is a convenient value—say 2^{15} or 2^{31} .

Efficient?

We can compute our hash values for s and the initial substring of t using compact evaluation.

$$p^{k-1} * c_1 + p^{k-2} * c_2 + \dots + p * c_{k-1} + c_k$$

- This requires a lot of multiplication!
- ▶ It can be re-written as...

$$h = c_k + p(c_{k-1} + p(c_{k-2} + \dots + p(c_3 + p(c_2 + pc_1))..))$$

- Where k = |s| and c_i is the i^{th} character of s.
- This requires |s| 1 multiplications and |s| 1 additions.

Efficiency!

- If we precompute $q = p^{k-1}$ we can find the next hash value, h' as:
- $h' = p * (h q * c_i) + c_j$ where we remove character i and add character j.
- \blacktriangleright This requires only 1 multiplication and 1 addition.
 - Constant time.
- Thus we have $\Theta(|s|)$ operations to perform the initial hashes and $|t| |s| * \Theta(1)$ operations to do the rehashing.
- ▶ Overall: our algorithm operates in $\Theta(|s| + |t|)$ time.
- We win!

Sorting

- We have seen few sorting algorithms
 - Insertion sort in worst-case takes $O(n^2)$
 - Merge Sort in worst-case takes O(nlogn)
 - Heap Sort in worst-case takes O(nlogn)
- Does this mean O(nlogn) is the minimum cost?
- Can we do better?

Two Models for Sorting

- Comparison-based sorting model
 - We will see that all comparison sort must take $\Omega(nlogn)$ comparisons.
- Other Sorts without relying on comparison
 - Bucket sort takes O(n) on average
 - Radix sort also works in linear time

Lower Bound for Comparison Sort

- Comparison Sort
 - Only use comparisons between elements to gain order
 - \circ E.g., given two inputs a_i and a_j , we perform one of comparisons to get the order between them
 - $a_i < a_j$
 - $a_i = a_j$
 - $a_i > a_j$
 - Examples include insertion sort, merge sort, heap sort, etc.

Examples-Insertion Sort

- Insertion sort uses the following strategy:
 - Start with the second element in the list.
 - Insert it in the right place in the preceding list.
 - Repeat with the next unsorted element.
 - Keep going until we have placed the last element in the list.



Examples-Merge Sort

Takes the strategy that recursively divides the unsorted array into two parts and merging them in order.

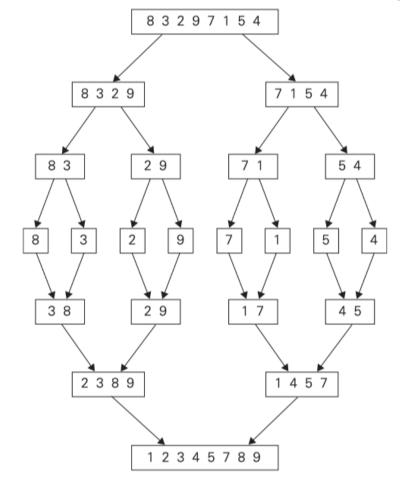


FIGURE 5.2 Example of mergesort operation.

Example-Heap Sort

```
Procedure heapsort (T[1..n])
       makeheap(T)
       for i = n to 2 step -1 do
              swap T[1] and T[i]
              siftdown (T[1 .. i - 1], 1)
              procedure siftdown (Heap, i)
              //move element i down to its correct position
                  c = i * 2
                  //Heap[c] < Heap[c+1] for a max-heap</pre>
                  if Heap[c] > Heap[c + 1]
                     c = c + 1
                  //for max-heap, the condition should be
                  //changed to Heap[i] < Heap[c]</pre>
                  if Heap[i] > Heap[c]
                     swap (Heap[i], Heap[c])
```

siftdown (Heap, c)

endif

end

A lower bound for the worst case

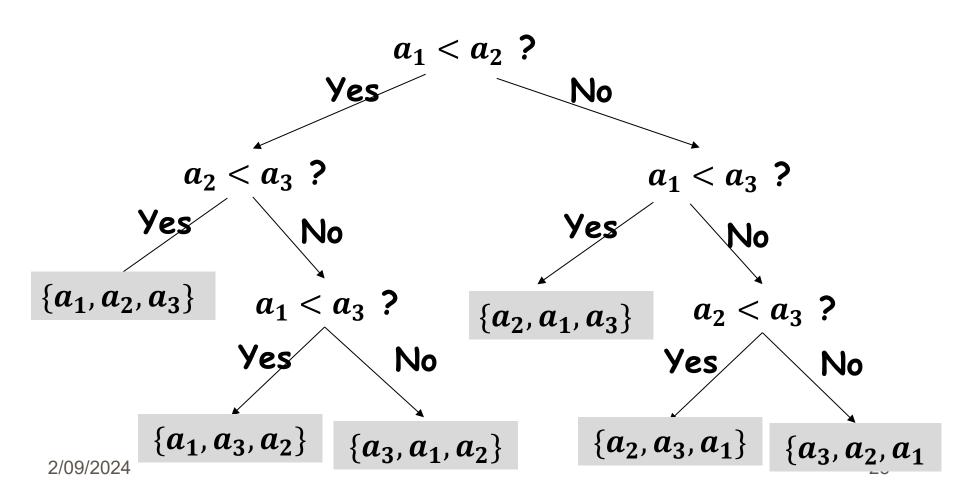
- Theorem:
 - All comparison sort must take $\Omega(nlogn)$ comparisons in the worst case.
- Next, we will show how to prove this theorem.

The Decision Tree Model

- A decision tree is a full binary tree that only represent the comparisons between elements.
- The comparisons order is determined by specific sorting algorithms.
- Control, data movement and all other inspect of algorithms are ignored.

Decision trees

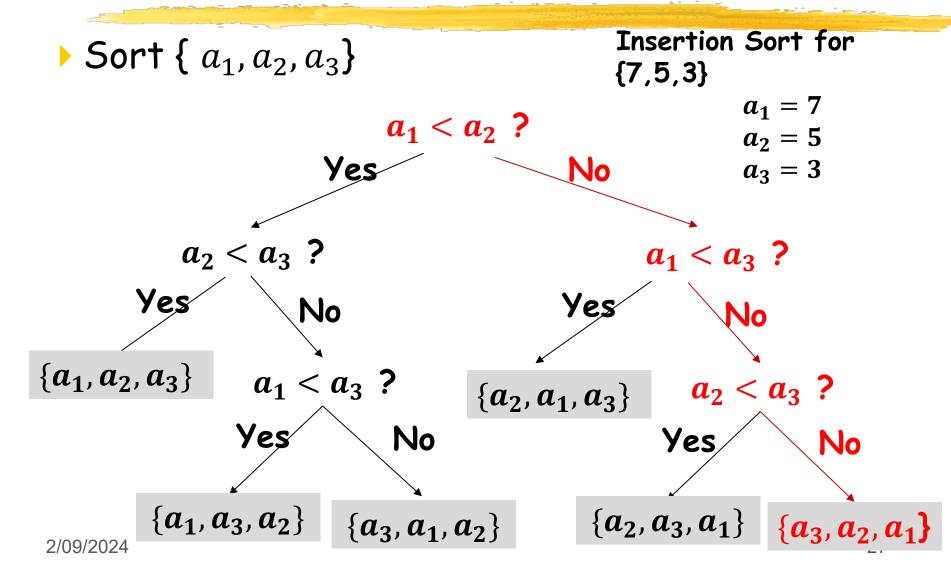
• Sort { a_1, a_2, a_3 }



Decision Tree

- \blacktriangleright Each internal node is a comparison of a_i and a_j
- All leaves nodes are all possible orderings of the items
- The execution of a sorting algorithm corresponds to tracing a simple path from the root of the decision tree down to a leaf.
- All comparison-based algorithms have an associated decision tree

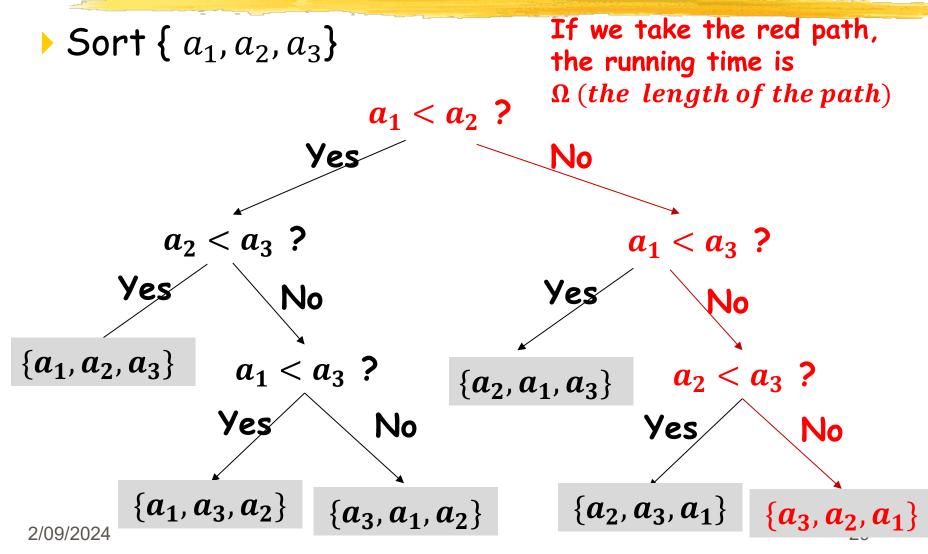
Decision trees



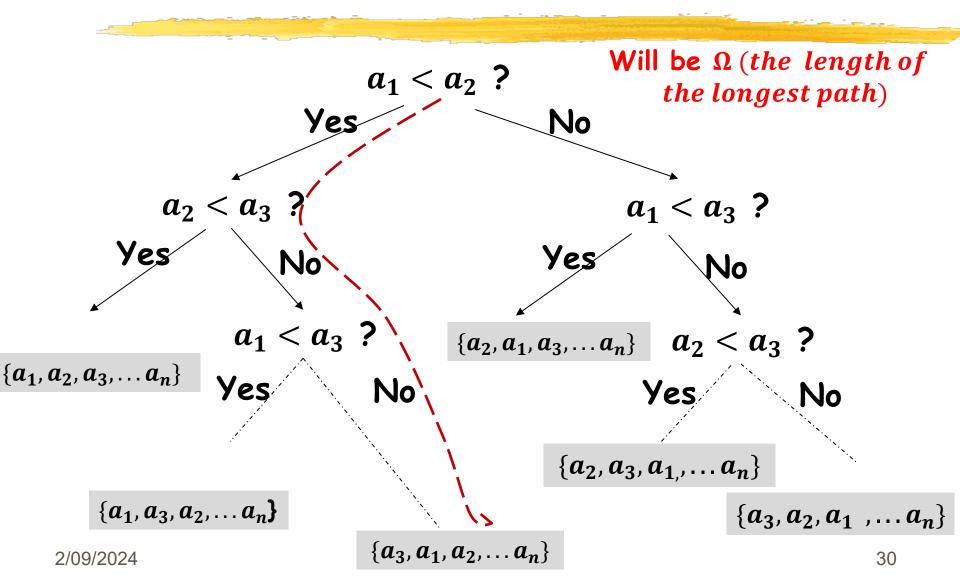
Decision Tree

- To sort n elements,
 - There will be n! permutations for n elements
 - There will be n! leaves in the decision tree
 - Any comparison algorithm must be able to produce each permutation of its input with size n, where each permutation will be a leaf node in the decision tree
 - Each of the leaf node is reachable from the root

What's the running time on a particular input?



What's the running time in the worst case?



How long is the longest path?

- The tree has n! leaves
- It is a binary tree
- We know that for a complete binary tree, the height will be $h = log_2^{(No.\,of\,\,leaves)}$
- The longest path is at least $log_2^{n!}$
- Vsing Stirling's formula $n! \sim \left(\frac{n}{e}\right)^n$
 - $\log_2^{n!} \sim \log_2\left(\frac{n}{e}\right)^n = n \log_2^{n/e} = \Omega(n\log n)$

Proof for the lower bound

- Any comparison sorting algorithm can be represented as a decision tree with n! leaves
- The worst running time is the longest length of path in that tree
- All decision tree with n! leaves have depth $\Omega(nlogn)$

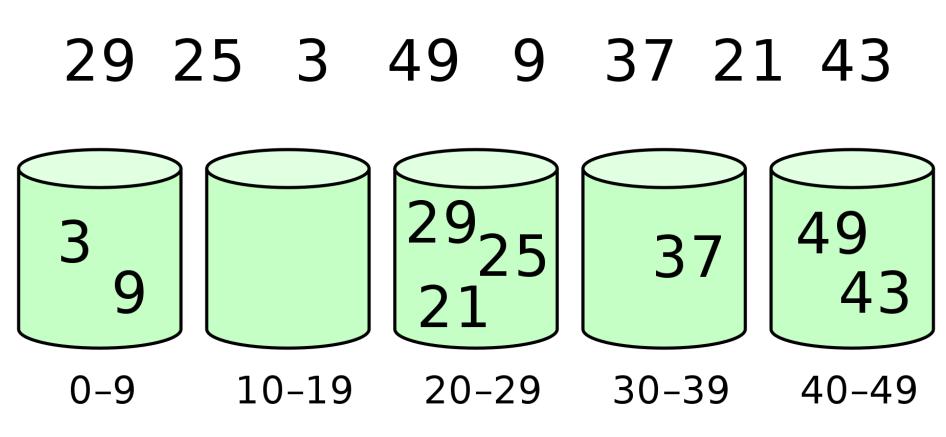
Corollary

- Heapsort and merge sort are asymptotically optimal comparison sorts.
 - The upper bound for these two types of sorting is O(nlogn), which matches the worst-case lower bound of comparison sorts.

Sorting in Linear Time

- Some sorting algorithm can run in linear time with specific requirement on its input.
- Bucket Sorting
 - Assume all inputs are drawn from uniform distribution with an average-case running time O(n)
 - works as follows
 - o divide input space into n buckets
 - distribute n inputs into the bucket

Bucket Sort



Sorted in O(n)

Bucket Sort

```
BUCKET-SORT(A)
   n = A.length
   let B[0..n-1] be a new array
                                                             → .39 /
   for i = 0 to n - 1
        make B[i] an empty list
                                                            → .68 /
    for i = 1 to n
                                                                   → .78 /
        insert A[i] into list B[\lfloor nA[i] \rfloor]
6
                                               10 .68
                                                             → .94 /
    for i = 0 to n - 1
                                                 (a)
                                                                  (b)
         sort list B[i] with insertion sort
    concatenate the lists B[0], B[1], \ldots, B[n-1] together in order
```

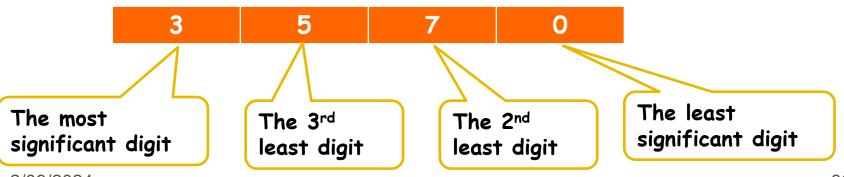
Each element is chosen from interval [0,1)

Bucket Sort

- We can analyze that the total running time of bucket sort is $\Theta(n)$
- We can implement each bucket as a linked list, which maintains the order for the element in the list
- Some issues:
 - Need to know the domain of elements to be sorted ahead of time

Radix Sort

- Works on decimal numbers
 - Each decimal digital has 10 possible values, 0,1,...,9
 - First sort the least significant digit
 - Then sort the 2nd,3rd until the most significant bits
 - It requires the sort to be stable



Radix Sort

Assume there are n elements in an array A each with d digits

```
RADIX-SORT(A, d)
```

- 1 for i = 1 to d
- 2 use a stable sort to sort array A on digit i

329		720		720		329
457	<u>)</u>]]]1-	355	jn-	329	jp.	355
657		436		436		436
839		457		839		457
436		657		355		657
720		329		457		720
355		839		657		839

Radix Sort

- To sort n d-digit numbers where each digit is chosen from k possible values, radix sort will take $\theta(d(n+k))$ time if the stable sort it uses takes $\theta(n+k)$ time.
- Proof Analysis
 - There are d iterations for sorting
 - Each iteration needs to sort n digits into k buckets, which is $\theta(n+k)$
- When d is constant and k=O(n), we make radix sort run in linear time

Related Reference

- Introduction to Algorithms, T. H. Cormen, 3rd Ed, MIT Press 2009.
 - Chapters 8
- Introduction to the Design and Analysis of Algorithms, A. Levitin, 3rd Ed., Pearson 2011.
 - Chapters 11.2